

Towards Structured Use of Bayesian Sequential Monitoring in Clinical Trials

Evan Kwiatkowski[†], Eugenio Andraca-Carrera[‡],
Mat Soukup[‡], Matthew A. Psioda^{†*}

[†] Department of Biostatistics, University of North Carolina,
McGavran-Greenberg Hall, CB#7420,
Chapel Hill, North Carolina, U.S.A.

[‡] Division of Biometrics VII, Office of Biostatistics
Center for Drug Evaluation and Research,
US Food and Drug Administration,
Silver Spring, Maryland, USA

August 30, 2019

Abstract

The text of your abstract. 200 or fewer words.

Keywords: 3 to 6 keywords, that do not appear in the title

*The authors gratefully acknowledge *please remember to list all relevant funding sources in the unblinded version*

1 Introduction

Things to discuss:

- 21st Century Cures Act (MATT)
- PDUFA VI reauthorization (MATT)
- Expansive work already done on sequential monitoring (EVAN – draft on 6/21)
- Our majors contribution (EVAN – as early as possible in introduction without having the flow appear weird – draft on 6/21)
- Outline for the remaining section of the paper (EVAN – draft on 6/21)

The theoretical foundations for the Bayesian clinical trials has been long established Cornfield (1966*a*) Cornfield (1966*b*) Neyman & Greenhouse (1967). These methods were not widely used in practice until a comprehensive framework for interpretation of results was developed through specifying prior distributions that were naturally and intuitively related to the research objectives (e.g. skeptical and enthusiastic priors) Freedman & Spiegelhalter (1989) Freedman & Spiegelhalter (1992) Spiegelhalter et al. (1993) Spiegelhalter et al. (1994) Fayers et al. (1997). (*Rewrite paragraph.*)

There is still potential for further utilization of Bayesian methods in the clinical trial setting. While the framework for interpretation of Bayesian clincial trials is well developed, the details of specifying prior distributions in a natural and intuitive way is lacking. This paper presents a structured or default way to determine prior distributions based on the trial design. Our major contribution is to present methods for the default or automatic selection of prior distributions in a way that is applicable to a wide array of clinical trial designs.

1. Bayesian methodology is widely developed.
2. It has been applied (cite).
3. The current perspective is that Bayesian methodology is only valid when Frequentist methods are insufficient, including where enrollment is challenging (rare diseases, pediatric studies)
4. Our contribution is to show that Bayesian methods are applicable to all clinical trials. This is shown by highlighting their improved interpretation and showing their use in varied and complicated situations.

2 Methods

As you introduce ideas that come from or extend other ideas in the literature, cite the relevant literature.

2.1 Monitoring versus Estimation Priors (EVAN – draft on 6/21)

2.1.1 Bayesian hypothesis testing based on posterior probabilities

The Bayesian paradigm provides direct inference on a parameter of interest through specification of a model for the data generating mechanism and prior distributions for unknown quantities. Let \mathbf{D} be a random variable representing the data collected in the trial with density $p(\mathbf{D}|\theta, \psi)$ where θ and ψ are the unknown quantities. Let θ be the parameter of interest and ψ be the unknown quantities that are not of primary importance (i.e. “nuisance parameters”). Define the sample spaces for the unknown quantities as $\theta \in \Theta$ and $\psi \in \Psi$.

Suppose the hypotheses under consideration are $H_0 : \theta \in \Theta_0$ versus $H_1 : \theta \in \Theta_1$. These hypotheses are adjudicated based on posterior probabilities of θ by evaluating its marginal likelihood

$$P(\theta \in \Theta_i | \mathbf{D}) = \int_{\Theta_i} p(\theta | \mathbf{D}) d\theta \text{ for } i \in \{0, 1\},$$

which is integrated or “marginalized” over the nuisance parameters $p(\theta | \mathbf{D}) = \int_{\Psi} p(\theta, \psi | \mathbf{D}) d\psi$.

2.1.2 Compelling level of evidence & prior elicitation

Define $\delta \in [0, 1]$ as a threshold for *a compelling level of evidence* as it relates to θ . We say that an individual is “all but convinced” that H_i is true given the observed data if $P(\theta \in \Theta_i | \mathbf{D}) \geq \delta$ for $i \in \{0, 1\}$. The quantity $1 - \delta$ reflects *residual uncertainty* of H_i being true relative to the competing hypothesis.

The posterior distribution of θ depends on the choice of prior distribution $\pi(\theta, \psi)$ since $p(\theta, \psi | \mathbf{D}) = p(\mathbf{D} | \theta, \psi) \pi(\theta, \psi) / p(\mathbf{D})$ by Bayes rule. The specification of the prior distribution depends on the research objective. An *inference prior* is a prior that is used when the research objective is to make final analysis after data collection is complete. A *monitoring prior* is a prior that is used when the research objective is to consider the impact of interim analyses on subject enrollment, with the potential for early termination of enrollment.

It has been said that “the purpose of a trial is to collect data that bring to conclusive consensus at termination opinions that had been diverse and indecisive at the *outset*” (Kass and Greenhouse (1989), emphasis added). These opinions manifest as priors $\pi(\theta, \psi)$ for which their relation to $P(\theta \in \Theta_i | \pi(\theta, \psi))$ $i \in \{0, 1\}$ is examined. Note this quantity does not depend on the data \mathbf{D} and therefore reflects a-priori opinion. A skeptical prior is an informative or subjective prior that gives substantial preference to H_0 such that it

is “all but convinced” that H_0 is true a-priori. This prior $\pi_S(\theta, \psi) \equiv \pi_S$ has the property that $P(\theta \in \Theta_0 | \pi_S) \geq \delta$ (equivalently $P(\theta \in \Theta_1 | \pi_S) < 1 - \delta$). The choice of $\delta \in [0, 1]$ is motivated by *a compelling level of evidence* as it relates to θ , although in this setting the “evidence” reflects a theoretical opinion rather than empirical judgement. For example, if $\delta = 0.95$, then this choice of skeptical prior places 95% prior probability that $\theta \in \Theta_0$. An enthusiastic prior $\pi_E(\theta, \psi) \equiv \pi_E$ similarly gives preference to H_1 through the property that $P(\theta \in \Theta_1 | \pi_E) \geq \delta$ (equivalently $P(\theta \in \Theta_0 | \pi_E) < 1 - \delta$). For purposes of interpretation, a *skeptical person* is someone whose a-priori opinions of θ are reflected through a skeptical prior and a *enthuastic person* is someone whose a-priori opinions of θ are reflected through an ethuastic prior. The prior distributions discussed are generally “non-informative” over the nuisance parameters.

2.1.3 Sequential monitoring & final inference

The use of monitoring based on changing the opinion of skeptical and enthuastic priors has been described as overcoming a handicap (Freedman & Spiegelhalter (1989)) and providing a brake (Fayers et al. (1997)) on the premature termination of trials, or constructing “an adversary who will need to be disilusioned by the data to stop further experimentation” (Spiegelhalter et al. (1994)).

Early termination of the trial is appriopriate if diverse prior opinions about θ would be in agreement given the interim data (e.g. the skeptical and enthuastic person reach the same conclusion). It is then reasonable to stop data collection if, upon seeing the data, a *skeptical person* changes their opinion to be “all but convinced” that H_1 is true ($P(\theta \in \Theta_1 | \mathbf{D}, \pi_S) \geq \delta$), or an *enthuastic person* becomes “all but convinced” that H_1 is false ($P(\theta \in \Theta_0 | \mathbf{D}, \pi_E) \geq \delta$).

The futility criteria isn't actually based on Θ_0 (e.g. it is based on an intermediate value in the proof-of-activity example). Therefore, develop the futility stopping criteria in its own section along with probability of success.

Final inference on θ is made once enrollment is stopped based on the monitoring priors or at the planned end of the trial. An inference prior $\pi_I(\theta, \psi) \equiv \pi_I$ is often non-informative or objective in the sense that it does not show a-priori preference to H_0 or H_1 ($P(\theta \in \Theta_0|\pi_I) \approx P(\theta \in \Theta_1|\pi_I)$). **Not always true. In the proof-of-activity example the inference prior which is a 50:50 mixture of skeptical and enthusiastic priors still favors H_1 .** There are many ways to formulate an inference prior with this property. We propose use of a mixture prior constructed from the monitoring process as the inference prior:

$$\pi_I = \omega \cdot \pi_S + (1 - \omega) \cdot \pi_E$$

for $\omega \in [0, 1]$. Choosing $\omega = 1/2$ for an equal mixture of π_S and π_E corresponds to an inference prior that is impartial to H_0 and H_1 , and is a practical choice of π_I is to be determined before the start of data collection. Define $p(\mathbf{D}|\pi(\theta, \psi)) = \int p(\mathbf{D}|\theta)\pi(\theta, \psi)d(\theta, \psi)$ to be the marginal likelihood for the data given the prior $\pi(\theta, \psi)$. Choosing ω based on posterior model probabilities of the null and alternative hypotheses yields $\omega = p(\mathbf{D}|\pi_S)/(p(\mathbf{D}|\pi_S) + p(\mathbf{D}|\pi_E))$.

The determination of a significant trial result is given by

$$P(\theta \in \Theta_1|\mathbf{D}, \pi_I) \geq \delta.$$

All relevant information about θ can be derived from its marginal posterior distribution with an inference prior (e.g. posterior mean, credible intervals). For example, the posterior

mean using the inference prior will be a two-part mixture of the posterior means using the skeptical and enthusiastic priors:

$$E(\theta|\mathbf{D}, \pi_I) = \omega \cdot E(\theta|\mathbf{D}, \pi_S) + (1 - \omega) \cdot E(\theta|\mathbf{D}, \pi_E)$$

$$\int_{\Theta} \theta p(\theta|\mathbf{D}, \pi_I) d\theta = \omega \cdot \int_{\Theta} \theta p(\theta|\mathbf{D}, \pi_S) d\theta + (1 - \omega) \cdot \int_{\Theta} \theta p(\theta|\mathbf{D}, \pi_E) d\theta$$

As an alternative strategy to futility analysis, one can monitor the probability of success (POS) for the trial. The probability of getting a convincing result at the end of the trail can be computed using the interim data. Let $p(\theta|\mathbf{D}, \pi_I)$ denote the posterior distribution for θ based on the inference prior π_I and the current data \mathbf{D} . Let ξ denote the POS which is given as follows:

$$\begin{aligned} \xi &= P[\mathbf{D}_1 \in \mathbb{R}^{dim(\mathbf{D}_1)} | P(\theta \in \Theta_1 | \mathbf{D}_1, \mathbf{D}, \pi_I) \geq \delta] \\ &= E[1\{P(\theta \in \Theta_1 | \mathbf{D}_1, \mathbf{D}, \pi_I) \geq \delta\}] \end{aligned}$$

where the expectation is taken with respect to the posterior predictive distribution $p(\mathbf{D}_1)$ for future data \mathbf{D}_1 (which includes subjects yet to enroll):

$$p(\mathbf{D}_1) = \int p(\mathbf{D}_1|\theta) \cdot \pi(\theta|\mathbf{D}) d\theta.$$

One may stop the enrollment if ξ is sufficiently small (i.e. $\xi < 0.05$).

3 Examples

3.1 Single-Arm Proof-of-Activity Trial with Binary Endpoint

3.1.1 Model formulation & prior elicitation

Consider a single-arm oncology proof-of-activity trial with a binary endpoint. The data \mathbf{D} are Binomially distributed and the response rate θ is the parameter of interest, with higher

values of θ being indicative of proof-of-activity.

Consider testing the hypothesis $H_0 : \theta \leq \theta_0$ versus $H_1 : \theta > \theta_0$. Consider a highly clinically relevant treatment effect $\theta_A > \theta_0$. A design property of the trial will be the posterior probability of $\theta \geq \theta_0$ given $\theta = \theta_A$ (similar to frequentist power).

Monitoring priors for this trial will be made using the concepts of a skeptical prior and an enthusiastic prior. Recall a skeptic is “all but convinced” that H_0 is true a-priori, therefore $P(\theta \in \Theta_0 | \pi_S) \geq \delta$. **(This expression of probability for the skeptic doesn’t hold in this example.)** An optimist is “all but convinced” that H_1 is true a-priori, therefore $P(\theta \in \Theta_1 | \pi_E) \geq \delta$.

It is intuitive to center the skeptical and enthusiastic priors around the quantities θ_0 and θ_A respectively, so that $E(\pi_S) = \theta_0$ and $E(\pi_E) = \theta_A$.

Beta priors for θ will be used to provide closed-form expressions of the posterior distributions via Beta-Binomial conjugacy (the posterior distribution $p(\theta | \mathbf{D})$ will be Beta distributed).

Consider the hypotheses

$$H_0 : \theta \leq 0.20$$

$$H_1 : \theta > 0.20$$

The skeptical prior will be Beta distributed, centered around $\theta_0 = 0.15$, and have 4.5% prior probability that $\theta > \theta_A$. Similarly, the enthusiastic prior will be centered around $\theta_A = 0.4$ and have 5% prior probability that $\theta < \theta_0$.

3.1.2 Sequential monitoring

The trial will proceed until one of the following three conditions are satisfied:

Efficacy criteria: $P(\theta > 0.20|\mathbf{D}, \pi_S) \geq 0.95$

Futility criteria: $P(\theta \leq 0.30|\mathbf{D}, \pi_E) \geq 0.85$

Maximum sample size: $N = 76$ patient outcomes obtained

Add justification for maximum sample size.

Let y_1 be the number of successes and y_0 be the number of failures.

Skeptical prior: $\pi_S(\theta) \sim \mathcal{B}(\alpha_S, \beta_S)$

Skeptical posterior: $p(\theta|\mathbf{D}, \pi_S) \sim \mathcal{B}(\alpha_S + y_1, \beta_S + y_0)$

Efficacy stopping: $P(\theta \in \Theta_1|\mathbf{D}, \pi_S) \geq \delta$

Enthusiastic prior: $\pi_E(\theta) \sim \mathcal{B}(\alpha_E, \beta_E)$

Enthusiastic posterior: $p(\theta|\mathbf{D}, \pi_E) \sim \mathcal{B}(\alpha_E + y_1, \beta_E + y_0)$

Futility stopping: $P(\theta \in \Theta_0|\mathbf{D}, \pi_E) \geq \delta$



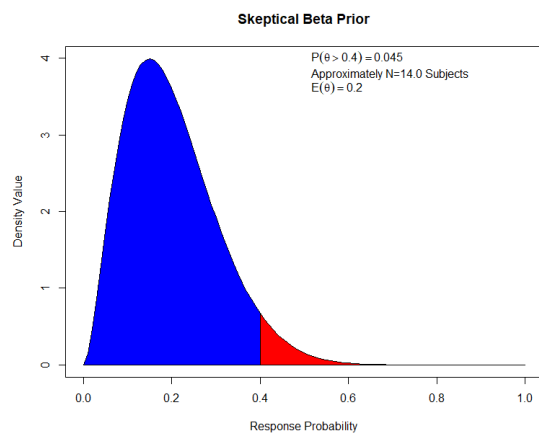


Figure 1: (a) Skeptical prior (b) Enthusiastic prior (c) 50/50 mixture of skeptical and enthusiastic prior

3.1.3 Example paths




Figure 2: (a) Early stopping for efficacy (b) Early stopping for futility

3.1.4 Design properties: Results

Assume that the outcomes are ascertained after approximately 4 months of follow-up and 2 patients per month on average are enrolled. The study may stop enrollment at a given point in time, but all enrolled patients will be followed for outcome ascertainment. An interim analysis will be completed after every 2 subjects complete follow-up.

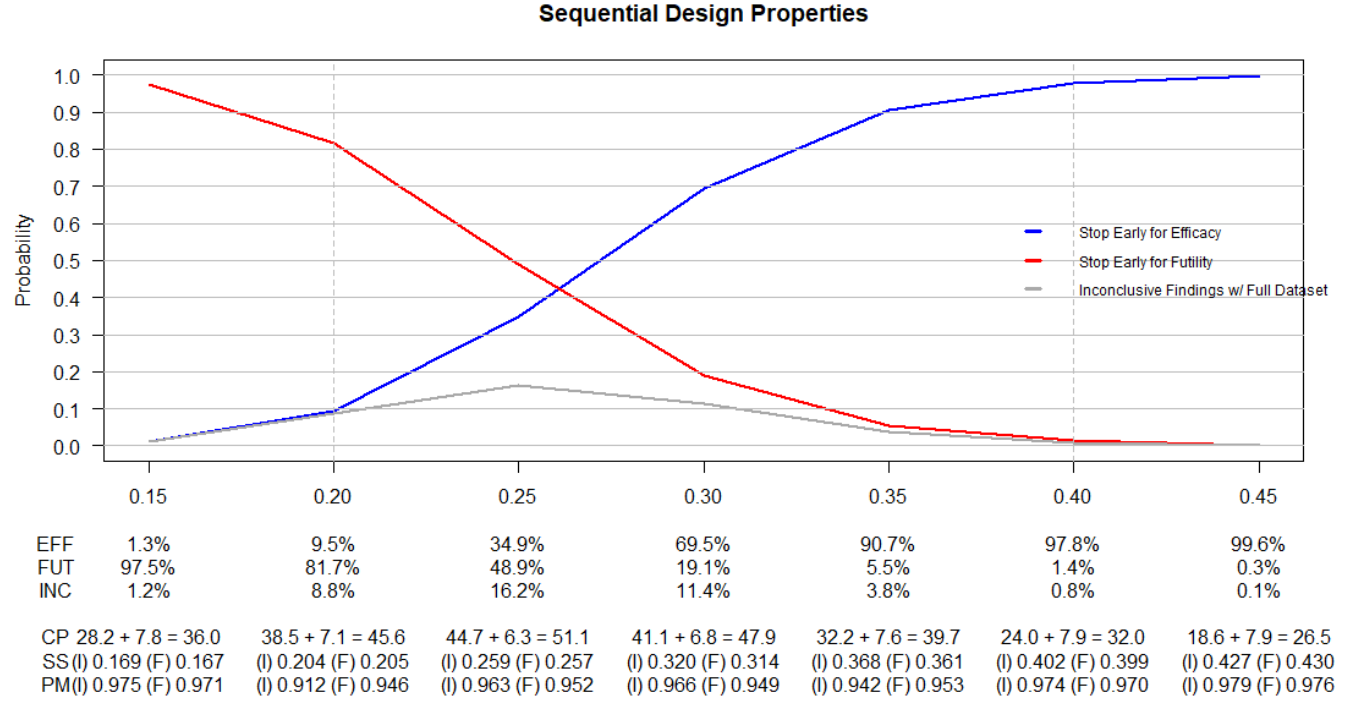


Figure 3: EFF is the probability of the trial stopping early for efficacy, FUT is the probability of the trial stopping early for futility, INC is the probability of reaching the maximum sample size without a conclusive monitoring result, SS is the average sample size at the interim analysis and at the end of follow-up (note there are approximately 8 subjects undergoing follow-up when enrollment is terminated), CP is the coverage probability using the mixture prior, PM is the posterior mean using the mixture prior.

3.1.5 Design properties: Slow vs. fast accrual

Add figure for how Type 1 error and power vary based on design schemes.

3.1.6 Agreement between interim and final result

- Use correct terminology: discuss the probability of stopping early vs. probability we conclude the treatment works.
- The two sides of the discussion: first is what happens during the trial regarding sequential monitoring, such as % of time stopping early vs. trial done to completion and expected sample size. Second is the final determination of efficacy or futility and how that relates to Type 1 Error and power.
- Make the point that regardless of frequency of monitoring there are good Type 1 error rates. Remember the best case for sequential monitoring is slow enrollment relative to outcome ascertainment. Slow enrollment means there is a benefit to ending trial early and reach a conclusion faster. Outcome ascertainment needs to be somewhat fast to ensure a good # of outcomes are generated.

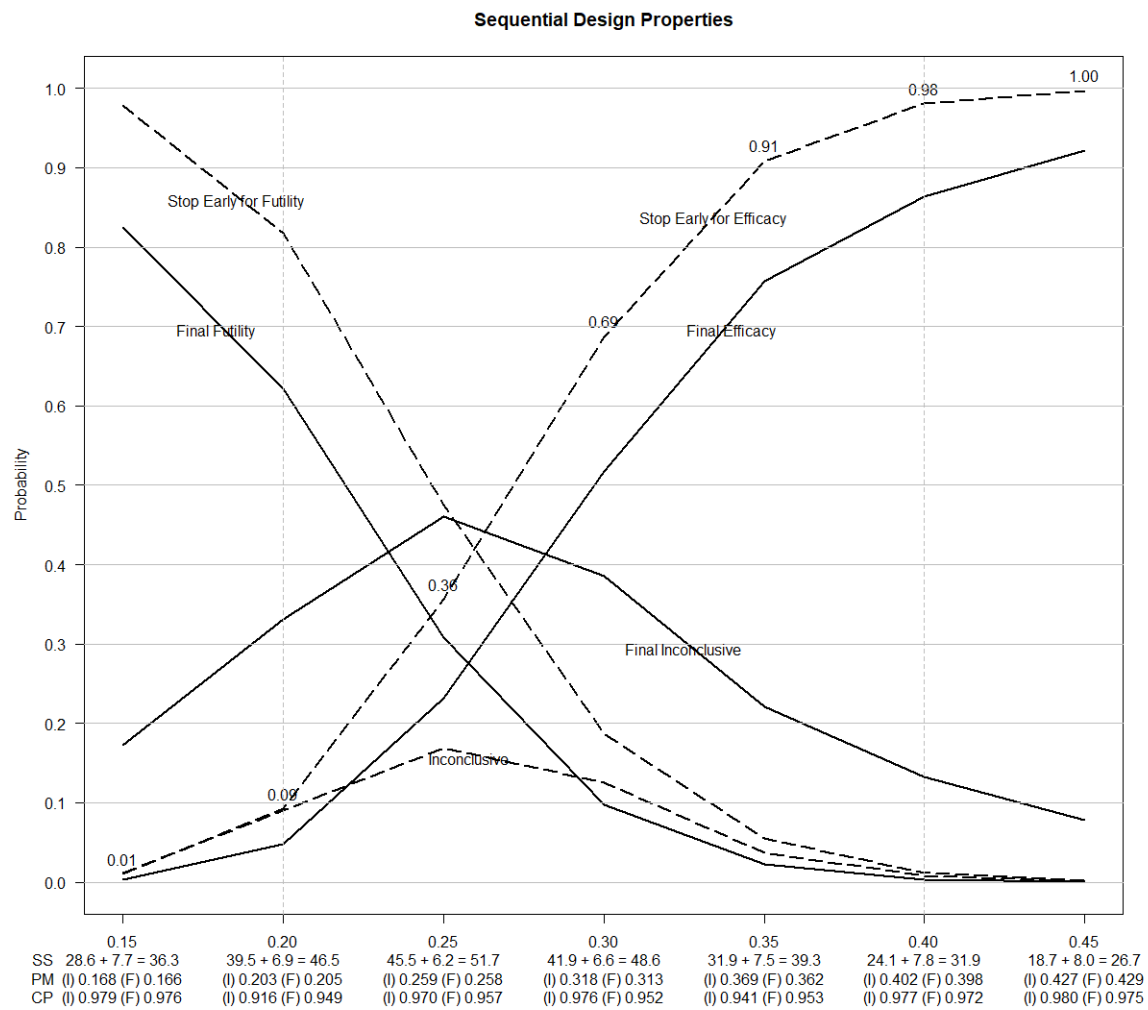


Figure 4: Shows agreement between sequential and final determinations as they relate to the strict cut-offs

3.1.7 Type 1 error rate by the frequency of data monitoring

- Want to highlight that the ultimate inference will have no Type 1 error inflation. At this point the ultimate inference for efficacy is still made with skeptical prior.
- Label lines nicely.
- “Only bad thing to do is to stop learning”
- Enrollment rate and % of ongoing data as operating characteristics are interesting ideas, but focus on the plots already created.
- Mat: Scaling on # of subjects in each interim analysis rather than # of interim analyses (e.g. flip X axis). Make labels go vertical or diagonal.

Bayesian monitoring has good frequentist properties even with frequent interim analyses.

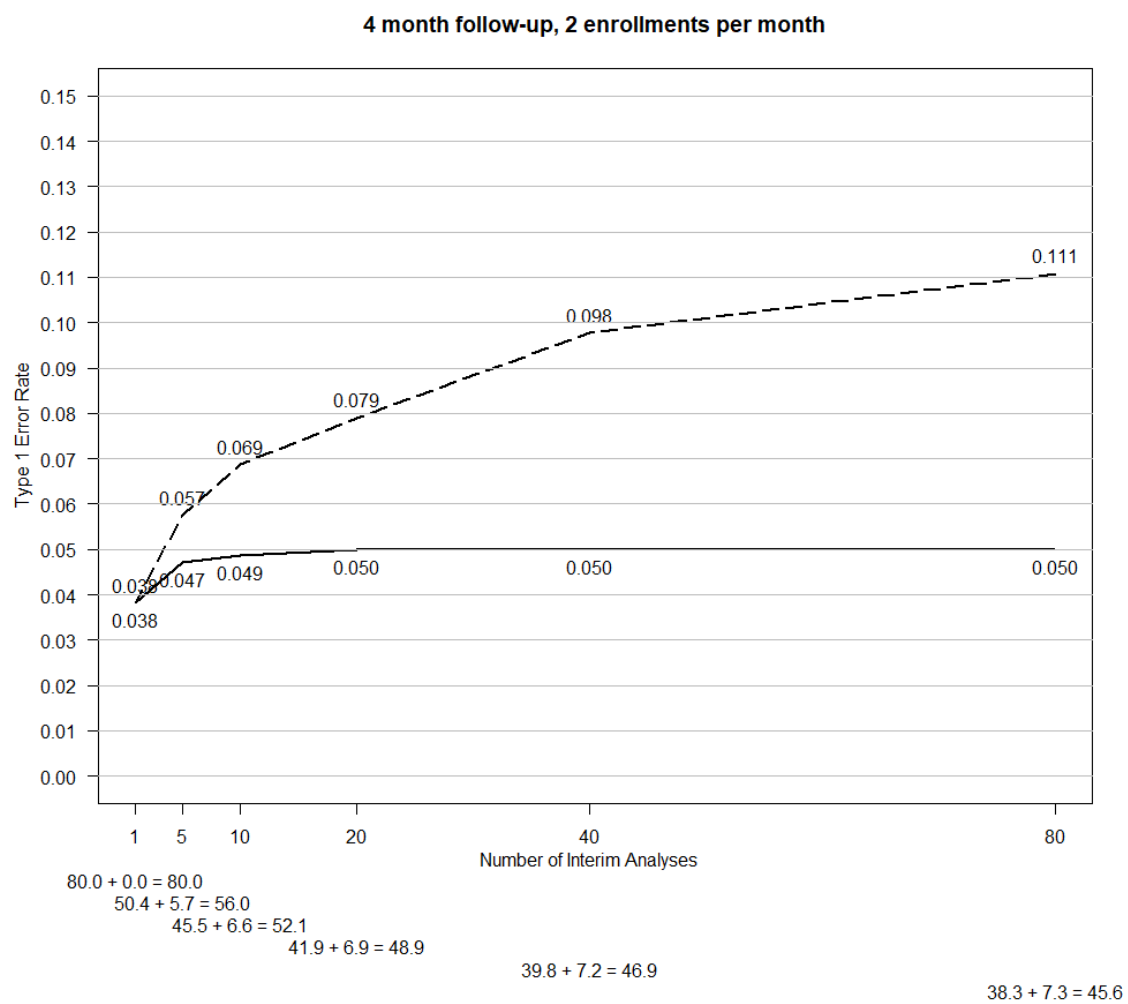



Figure 5: Type 1 error rate depending on frequency of sequential monitoring

Sequential Design Properties – Slow Accrual +2/month



Note: % AGR. = Percent of agreement between final and interim posterior probabilities relative to 0.95 threshold.

Figure 6: Have the necessary data summaries in R, just need to replicate the plot.

3.1.8 Evidence decrease

- Plot is conditional on evidence going down. For example, at $\theta = 0.45$ there is a 93% chance of agreement, which means there is evidence decrease only 7% of the time. At $\theta = 0.15$ there is evidence decrease 64% of the time. Note that minimum is very sensitive to the number of simulations and should be removed from the plot.
- Note that there is functionally a minimum sample size already that is the # of subjects needed to overwhelm the skeptical prior. Actually compute the # (perhaps after 5

consecutive successes).

- There is a drop in the final power which is undesirable, but the drop in type 1 error is desirable. These move in tandem.
- Move discussion away from being scared of evidence decrease. There is an attenuation with respect to dichotomous threshold, but not overall. In this sense % agreement may be misleading. Look at the % of change in the evidence rather than the % change in the agreement based on a dichotomous agreement. The methods section may need to be modified to reduce focus on dichotmous decision. However, don't abandon hypothesis testing, but look at it holistically. Offer different perspectives on attenuation in the discussion. After all, 0.93 posterior probability isn't that different from 0.95.

4 Robustness of parameterizations of monitoring priors

- Test the hypothesis that power curves will be identical and what will change is the expected sample sizes and that varying the decay rate will affect bias not Type I/II error
- Actually, the having one prior “default” and the other spike-slab changes the monitoring result substantially. When using the spike-slab skeptical prior and the default enthusiastic prior, the determination of efficacy at the interim is made less frequently than when using the default skeptical prior. Similarly, when using the spike-slab enthusiastic prior and the default skeptical prior, the determination of futility is made less frequently than when using the default enthusiastic prior.
- The spike-slab and flattened priors are both 2-part mixtures of Beta distributions. The flattened priors are not very well represented by 2-part mixtures of Betas. I would have to add more Beta components to produce good representations of a flattened prior.




Figure 7: (a) Spike-slab skeptical prior (b) Flattened skeptical prior (c) Spike-slab enthusiastic prior (d) Flattened enthusiastic prior

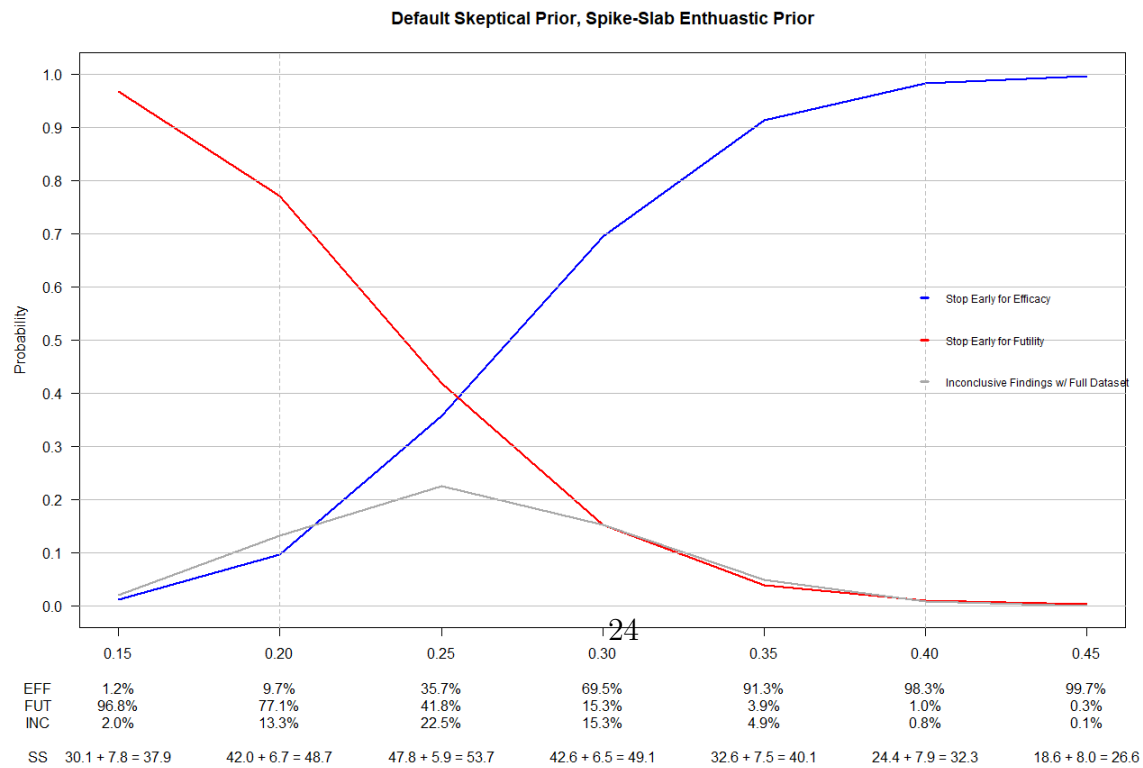
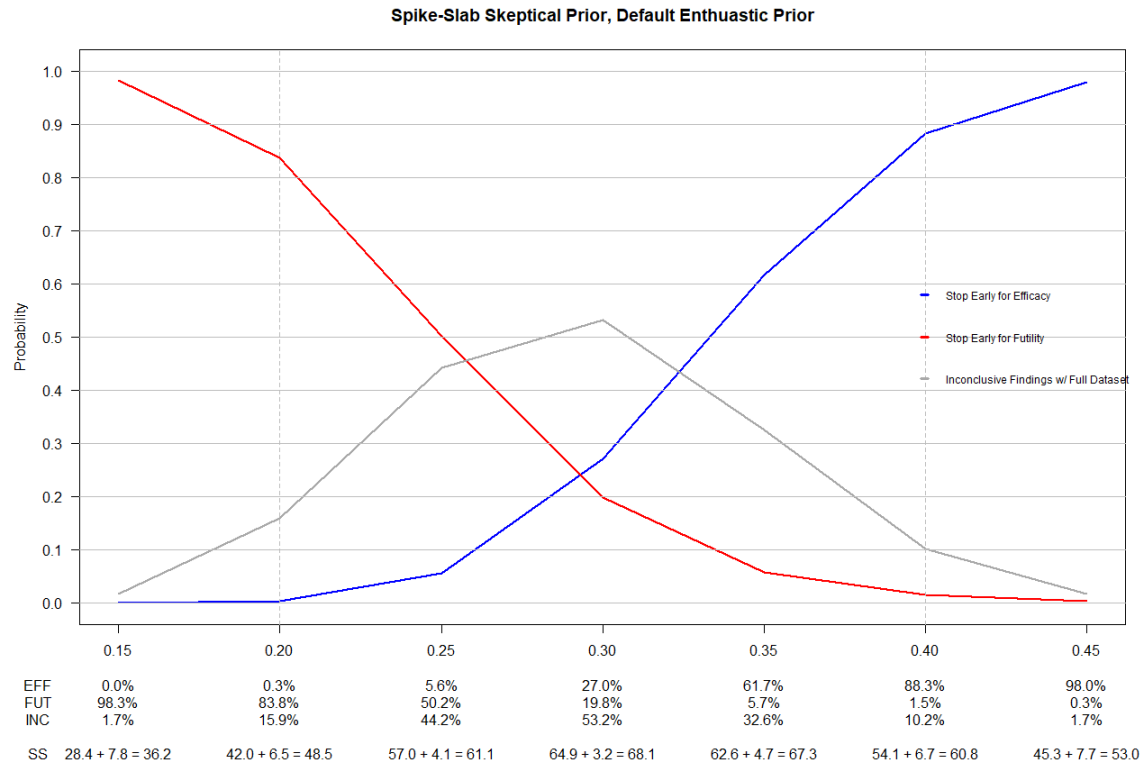


Figure 8: (a) Spike-slab skeptical prior and default enthusiastic prior (b) Default skeptical prior and spike-slab enthusiastic prior

4.1 Parallel Two-Group Superiority Trial /w Continuous Binary Endpoint

Interesting because prior is on risk difference $[-1,1]$ while also being non-informative on control group. Will need numerical integration to evaluate posteriors.

4.2 Three-Arm, Placebo Controlled Non-Inferiority Trial w/ Continuous Endpoint

$$P \rightarrow \beta_0 \text{ (placebo)}$$

$$C \rightarrow \beta_0 + \beta_1 \text{ (control)}$$

$$A \rightarrow \beta_0 + \beta_1 + \beta_2 \text{ (active)}$$

$$H_0 : \beta_2 - \delta\beta_1 \leq 0$$

Parameters of interest (β_1, β_2) , nuisance parameters (β_0, σ^2) .

Need priors $\pi(\beta_0), \pi(\beta_1), \pi(\beta_2|\beta_1)$.

Will use MCMC to evaluate posteriors.

5 Discussion – (MATT/EVAN)

Q: Why not reverse engineer priors to have exact Type 1 error properties?

A: This would basically be a frequentist method, in that the design would have to be adhered to exactly (including number and timing of data monitoring). Philosophically, designing a Bayesian study that requires rigid monitoring rules loses the advantages of Bayes from the likelihood principle.

SUPPLEMENTARY MATERIAL

6 Beta Priors

Beta priors for θ will be used to provide closed-form expressions of the posterior distributions via Beta-Binomial conjugacy (the posterior distribution $p(\theta|\mathbf{D})$ will be Beta distributed). The Beta distribution has two shape parameters. These parameters can be determined uniquely by specifying the desired mean and variance of the distribution. The variance for the skeptical and enthusiastic priors is then uniquely determined through by the choice of threshold δ . In particular, let $\pi_S(\theta) \sim \mathcal{B}(\alpha, \beta)$ be Beta distributed with shape parameters (α, β) . There is a single choice of (α, β) such that:

$$\theta_0 = E(\pi_S) = \int_{\Theta} \pi_S(\theta) d\theta = \frac{\alpha}{\alpha + \beta} \text{ and } \delta = \int_{\Theta_0} \pi_S(\theta) d\theta = \int_0^{\theta_0} \frac{\theta^{\alpha-1}(1-\theta)^{\beta-1}}{B(\alpha, \beta)} d\theta$$

where $B(\alpha, \beta)$ is the Beta function.

Alternatively, the variance could be determined by specifying a desired quantile of the prior distribution which would then be reflected in δ . Then there is a single choice of (α, β) such that

$$\theta_0 = E(\pi_S) = \int_{\Theta} \pi_S(\theta) d\theta = \frac{\alpha}{\alpha + \beta} \text{ and } \lambda = \int_{\theta_A}^1 \pi_S(\theta) d\theta = \int_{\theta_A}^1 \frac{\theta^{\alpha-1}(1-\theta)^{\beta-1}}{B(\alpha, \beta)} d\theta,$$

in which case $\delta = \int_{\Theta_0} \pi_S(\theta) d\theta$ is a deterministic quantity.

7 BibTeX

References

Cornfield, J. (1966*a*), ‘A Bayesian Test of Some Classical Hypotheses, with Applications to Sequential Clinical Trials’, *Journal of the American Statistical Association* **61**(315), 577.
URL: <https://www.jstor.org/stable/2282772?origin=crossref>

Cornfield, J. (1966*b*), ‘Sequential Trials, Sequential Analysis and the Likelihood Principle’, *The American Statistician* **20**(2), 18.
URL: <https://www.jstor.org/stable/2682711?origin=crossref>

Fayers, P. M., Ashby, D. & Parmar, M. K. B. (1997), ‘Tutorial in Biostatistics: Bayesian Data Monitoring in Clinical Trials’, *Statistics in Medicine* **16**(12), 1413–1430.
URL: <http://doi.wiley.com/10.1002/%28SICI%291097-0258%2819970630%2916%3A12%3C1413%3A%3E3.0.CO%3B2-U>

Freedman, L. S. & Spiegelhalter, D. J. (1989), ‘Comparison of Bayesian with group sequential methods for monitoring clinical trials’, *Controlled Clinical Trials* **10**(4), 357–367.
URL: <https://www.sciencedirect.com/science/article/pii/0197245689900019?via%3Dihub>

Freedman, L. S. & Spiegelhalter, D. J. (1992), ‘Application of bayesian statistics to decision making during a clinical trial’, *Statistics in Medicine* **11**(1), 23–35.
URL: <http://doi.wiley.com/10.1002/sim.4780110105>

Hyman, D. M., Puzanov, I., Subbiah, V., Faris, J. E., Chau, I., Blay, J.-Y., Wolf, J., Raje, N. S., Diamond, E. L., Hollebecque, A., Gervais, R., Elez-Fernandez, M. E., Italiano, A., Hofheinz, R.-D., Hidalgo, M., Chan, E., Schuler, M., Lasserre, S. F., Makrutzki, M.,

Sirzen, F., Veronese, M. L., Tabernero, J. & Baselga, J. (2015), ‘Vemurafenib in Multiple Nonmelanoma Cancers with *BRAF* V600 Mutations’, *New England Journal of Medicine* **373**(8), 726–736.

URL: <http://www.nejm.org/doi/10.1056/NEJMoa1502309>

Neyman, J. & Greenhouse, S. W. (1967), *Proceedings of the Berkeley Symposium on Mathematical Statistics and Probability.*, University of California Press.

URL: <https://projecteuclid.org/euclid.bsmsp/1200513830>

Spiegelhalter, D. J., Freedman, L. S. & Parmar, M. K. B. (1993), ‘Applying Bayesian ideas in drug development and clinical trials’, *Statistics in Medicine* **12**(15-16), 1501–1511.

URL: <http://doi.wiley.com/10.1002/sim.4780121516>

Spiegelhalter, D. J., Freedman, L. S. & Parmar, M. K. B. (1994), ‘Bayesian Approaches to Randomized Trials’, *Journal of the Royal Statistical Society. Series A (Statistics in Society)* **157**(3), 357.

URL: <https://www.jstor.org/stable/10.2307/2983527?origin=crossref>